

Chapter 3

Transport Layer

A note on the use of these PowerPoint slides:


We're making these slides freely available to all (faculty, students, readers). They're in PowerPoint form so you see the animations; and can add, modify, and delete slides (including this one) and slide content to suit your needs. They obviously represent a *lot* of work on our part. In return for use, we only ask the following:

- If you use these slides (e.g., in a class) that you mention their source (after all, we'd like people to use our book!)
- If you post any slides on a www site, that you note that they are adapted from (or perhaps identical to) our slides, and note our copyright of this material.

For a revision history, see the slide note for this page.

Thanks and enjoy! JFK/KWR

All material copyright 1996-2020
J.F Kurose and K.W. Ross, All Rights Reserved



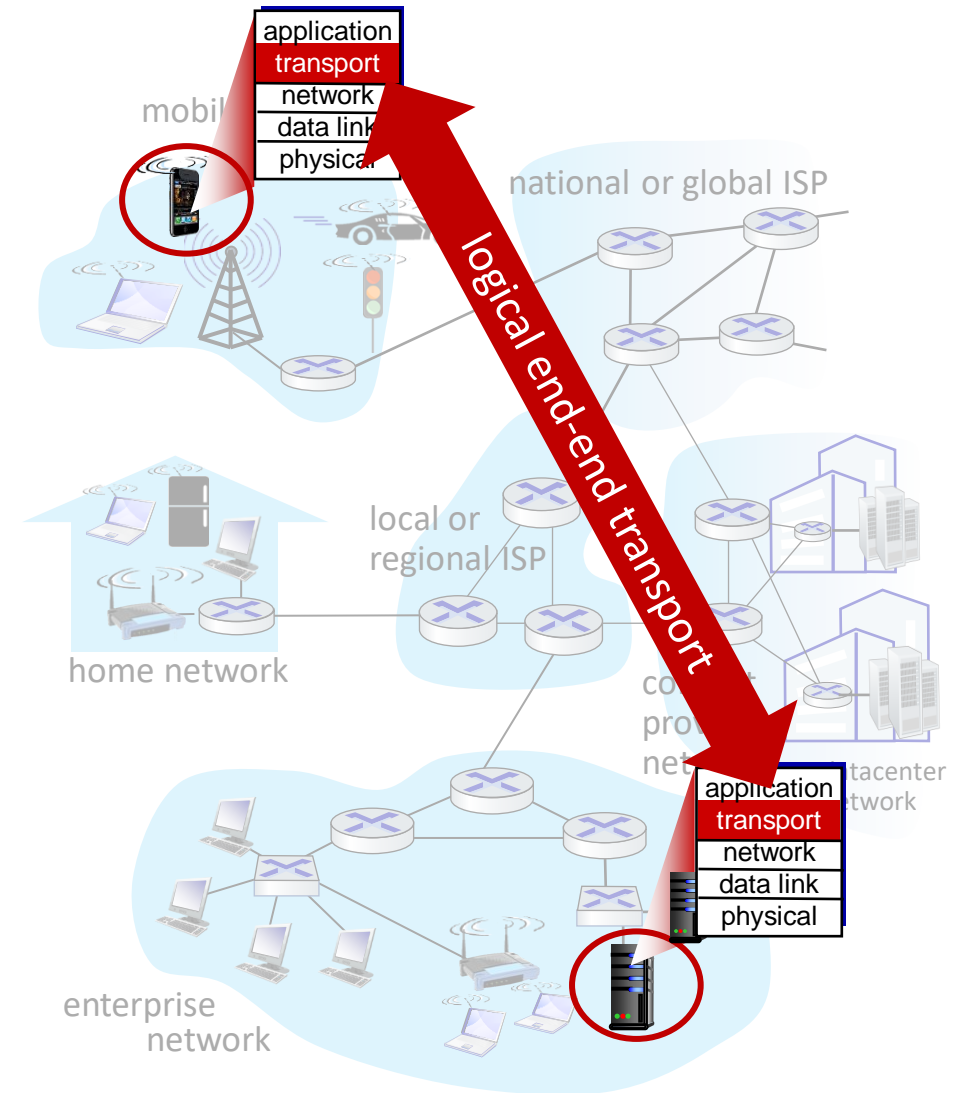
Computer Networking: A Top-Down Approach

8th edition

Jim Kurose, Keith Ross
Pearson, 2020

Transport services and protocols

- provide *logical communication* between application processes running on different hosts
- transport protocols actions in end systems:
 - sender: breaks application messages into *segments*, passes to network layer
 - receiver: reassembles segments into messages, passes to application layer
- two major transport protocols available to Internet applications
 - TCP, UDP



Transport vs. network layer services and protocols

- **network layer:** logical communication between *hosts*
- **transport layer:** logical communication between *processes*
 - relies on, enhances, network layer services

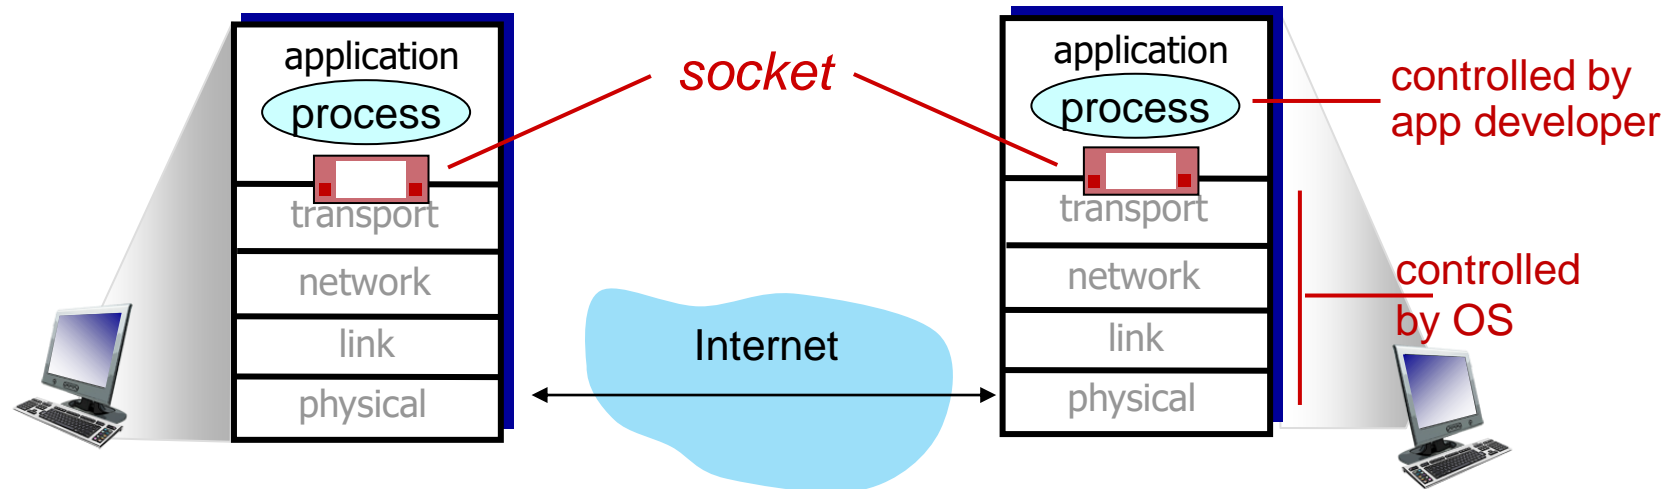
household analogy:

12 kids in Ann's house sending letters to 12 kids in Bill's house:

- hosts = houses
- processes = kids
- app messages = letters in envelopes
- transport protocol = Ann and Bill
- network-layer protocol = postal service

A fundamental concept: the socket

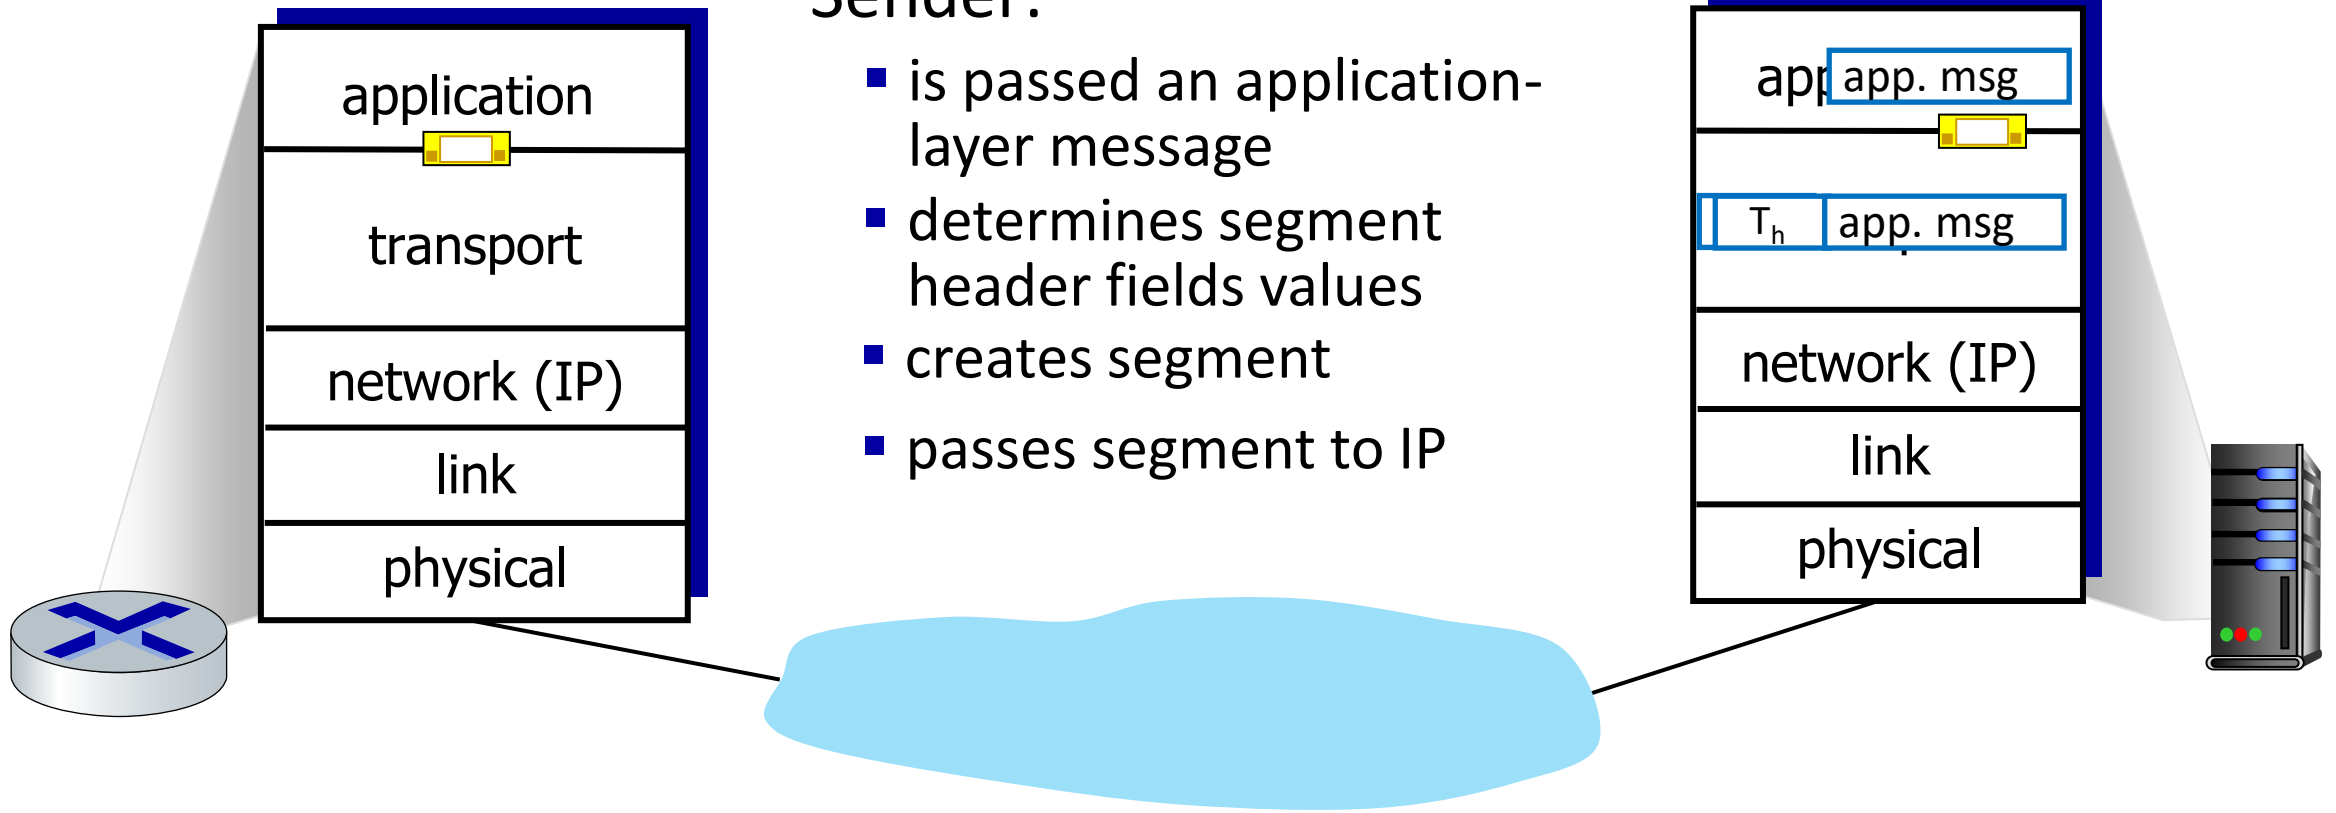
- process sends/receives messages to/from its **socket**
- socket analogous to door
 - sending process shoves message outdoor
 - sending process relies on transport infrastructure on other side of door to deliver message to socket at receiving process
 - two sockets involved: one on each side



Transport Layer Actions

Sender:

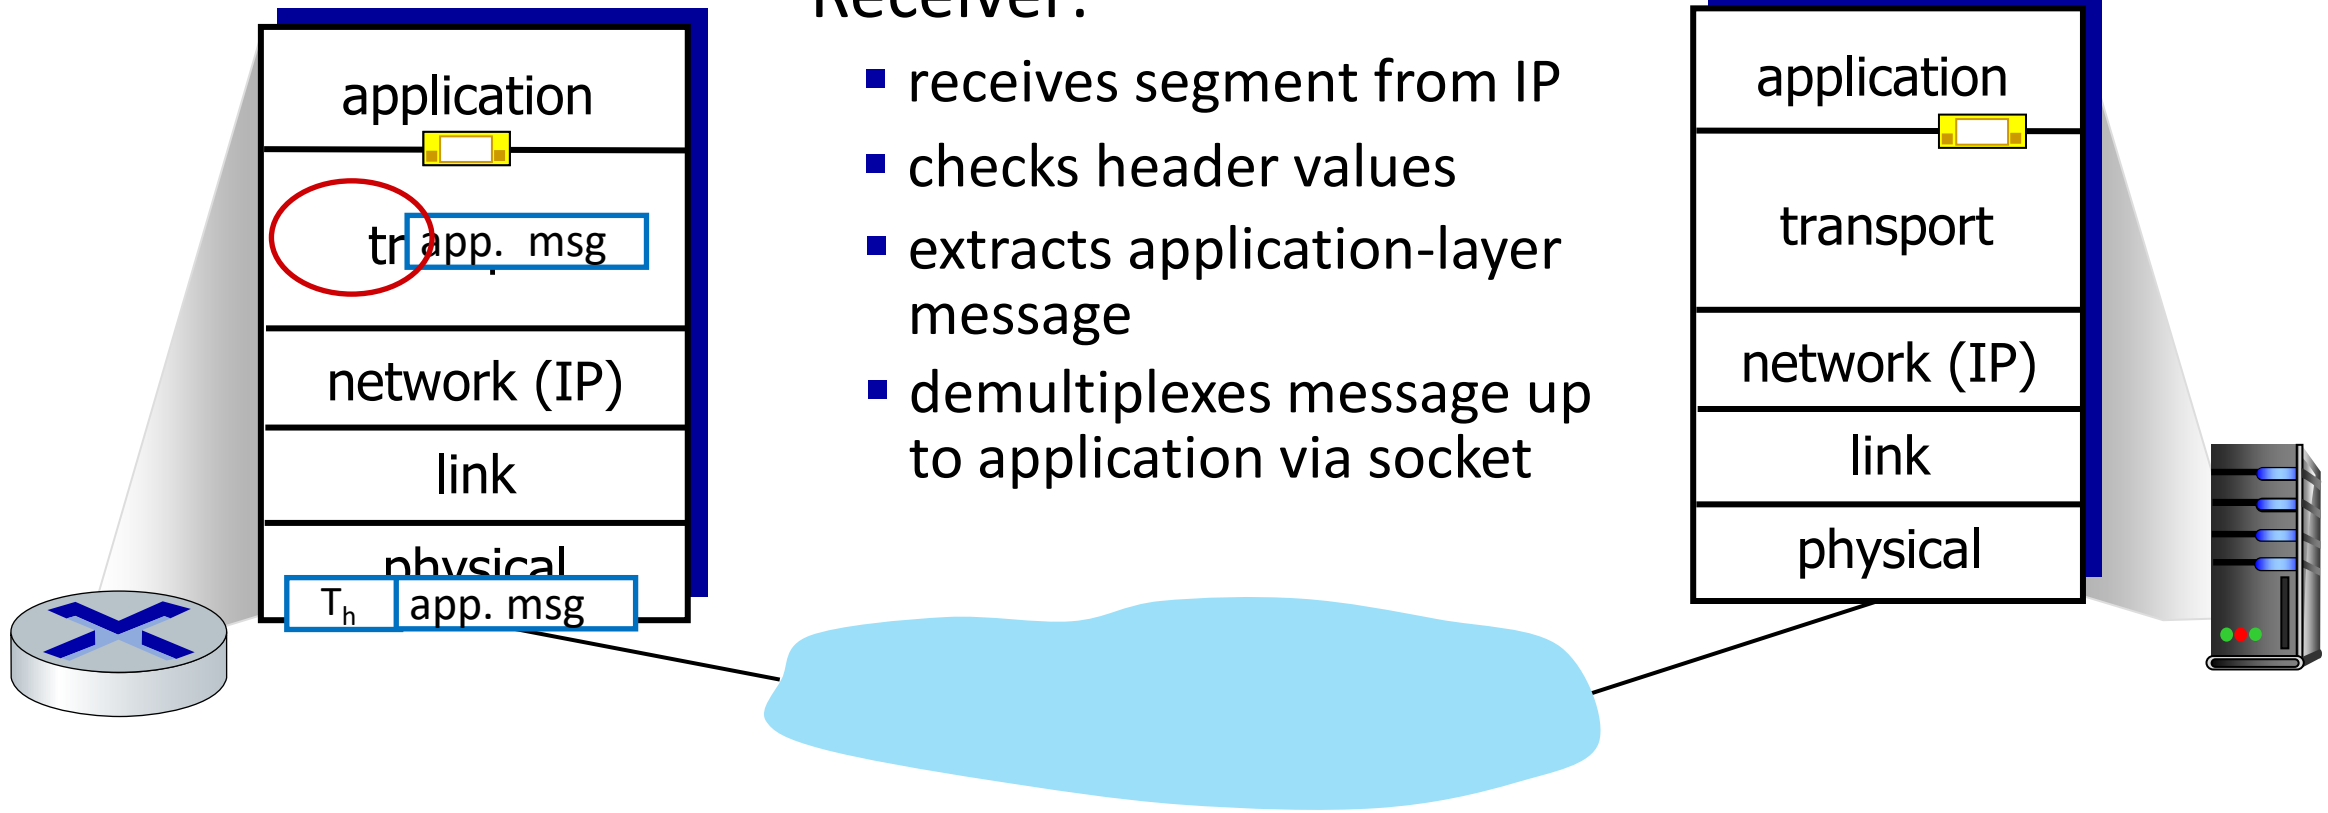
- is passed an application-layer message
- determines segment header fields values
- creates segment
- passes segment to IP



Transport Layer Actions

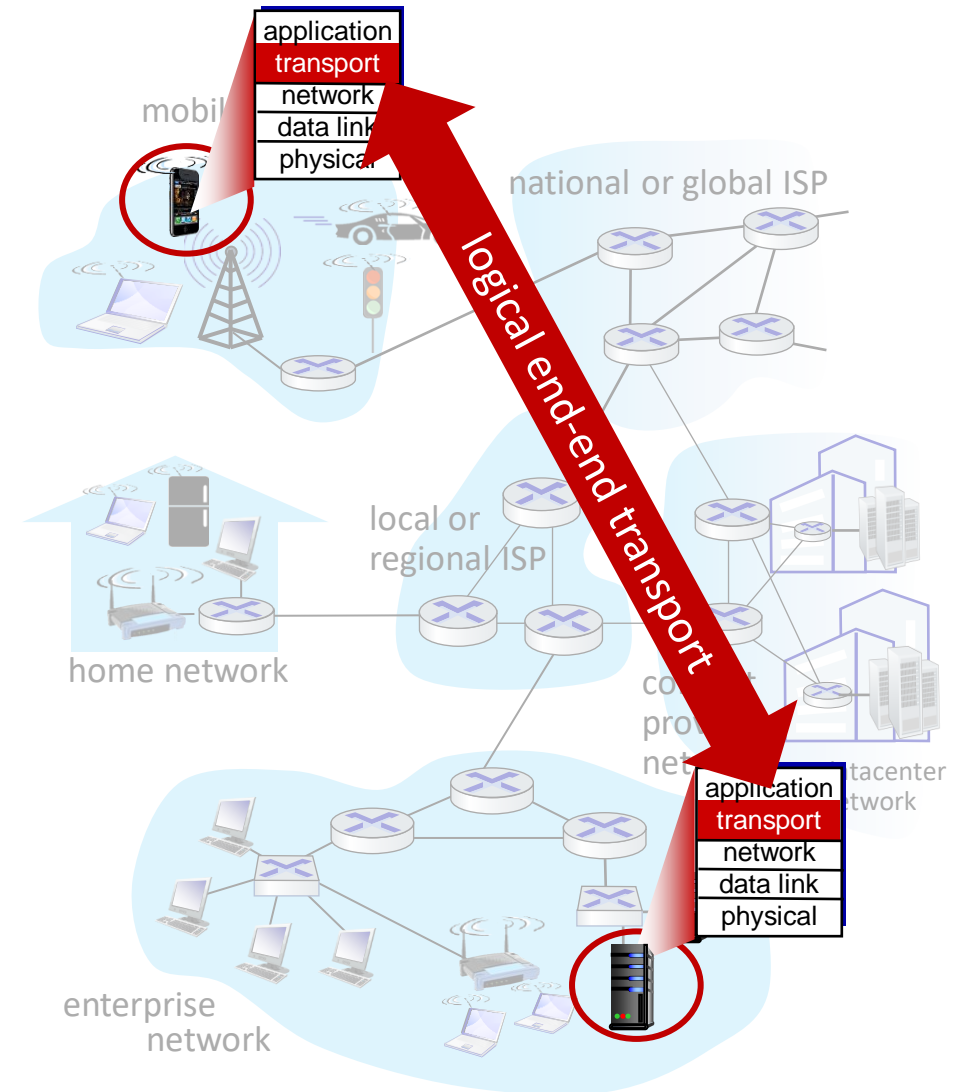
Receiver:

- receives segment from IP
- checks header values
- extracts application-layer message
- demultiplexes message up to application via socket



Two principal Internet transport protocols

- **TCP:** Transmission Control Protocol
 - reliable, in-order delivery
 - congestion control
 - flow control
 - connection setup
- **UDP:** User Datagram Protocol
 - unreliable, unordered delivery
 - no-frills extension of “best-effort” IP
- services not available:
 - delay guarantees
 - bandwidth guarantees



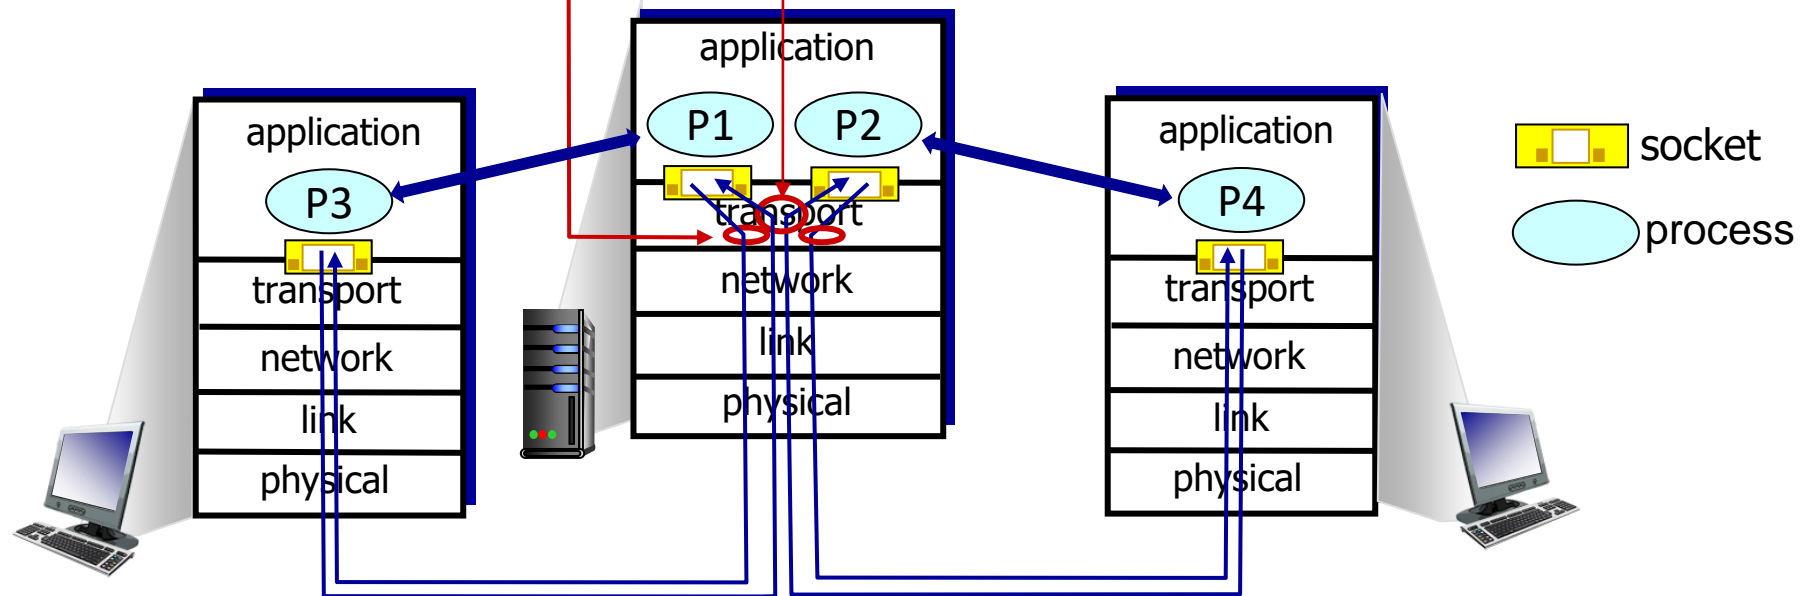
Multiplexing/demultiplexing

multiplexing at sender:

handle data from multiple sockets, add transport header (later used for demultiplexing)

demultiplexing at receiver:

use header info to deliver received segments to correct socket



How demultiplexing works

- host receives IP datagrams
 - each datagram has source IP address, destination IP address
 - each datagram carries one transport-layer segment
 - each segment has source, destination port number
- host uses *IP addresses & port numbers* to direct segment to appropriate socket

Connectionless demultiplexing

Recall:

- when creating socket, must specify *host-local* port #:

```
DatagramSocket mySocket1  
= new DatagramSocket(12534);
```

- when creating datagram to send into UDP socket, must specify

1. *destination IP address*
2. *destination port #*

- these two values univocally identify a UDP socket!

when receiving host receives *UDP* segment:

- checks destination port # in segment
- directs UDP segment to socket with that port #



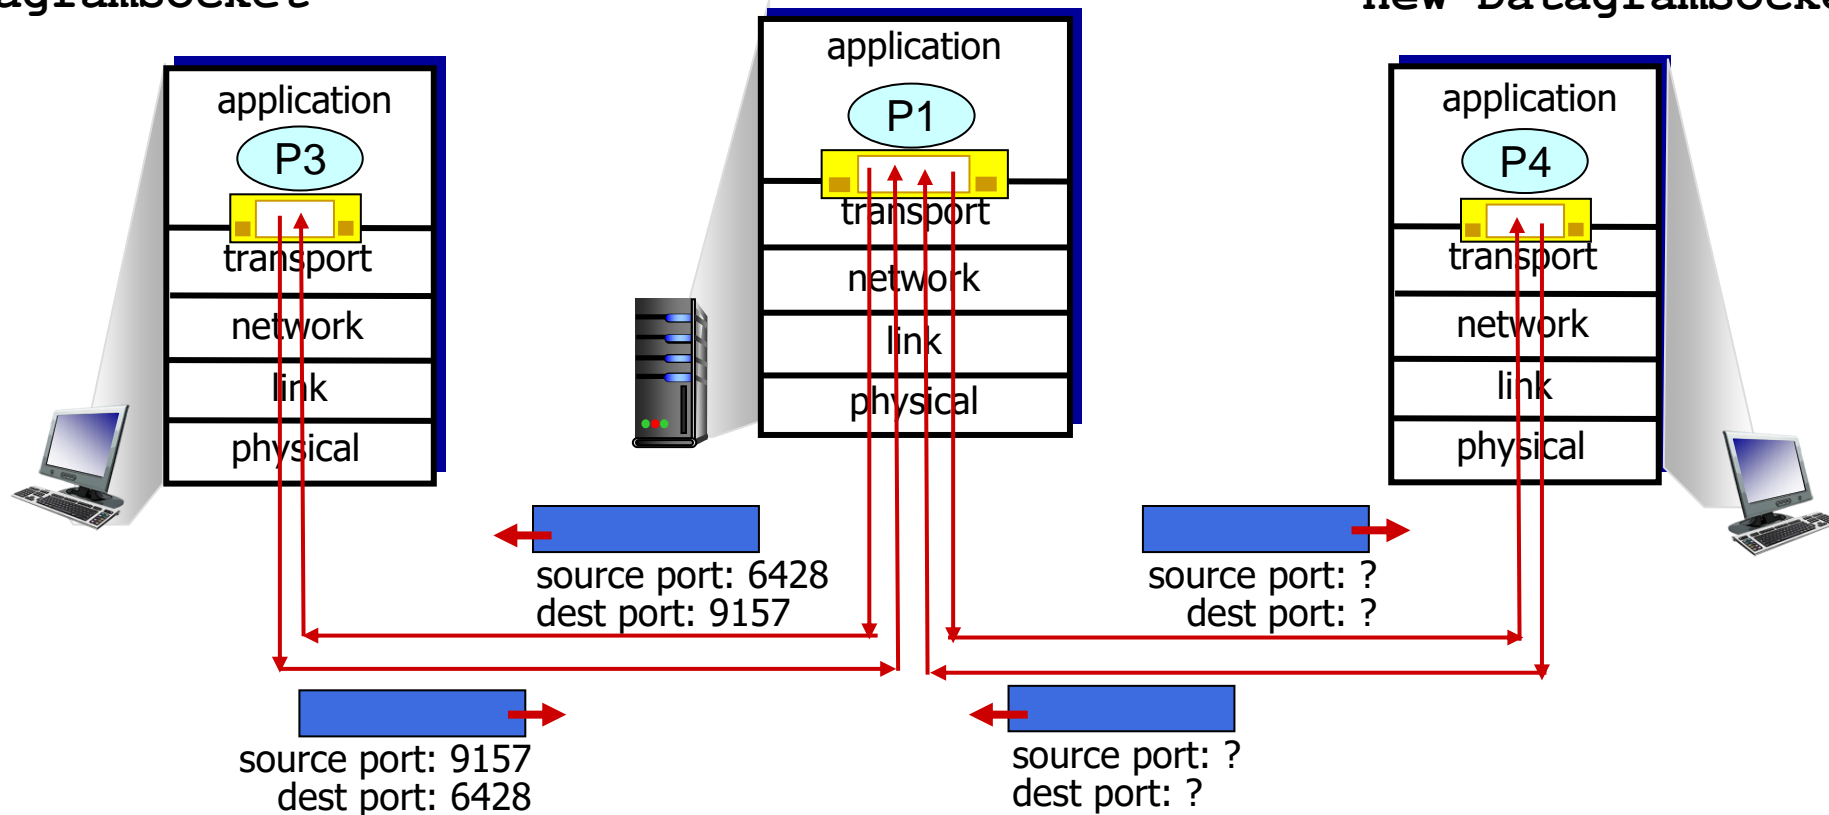
IP/UDP datagrams with *same dest. port # and dest. IP addresses*, but different source IP addresses and/or source port numbers will be directed to *same socket* at receiving host

Connectionless demultiplexing: an example

```
DatagramSocket mySocket2 =  
new DatagramSocket  
(9157);
```

```
DatagramSocket  
serverSocket = new  
DatagramSocket  
(6428);
```

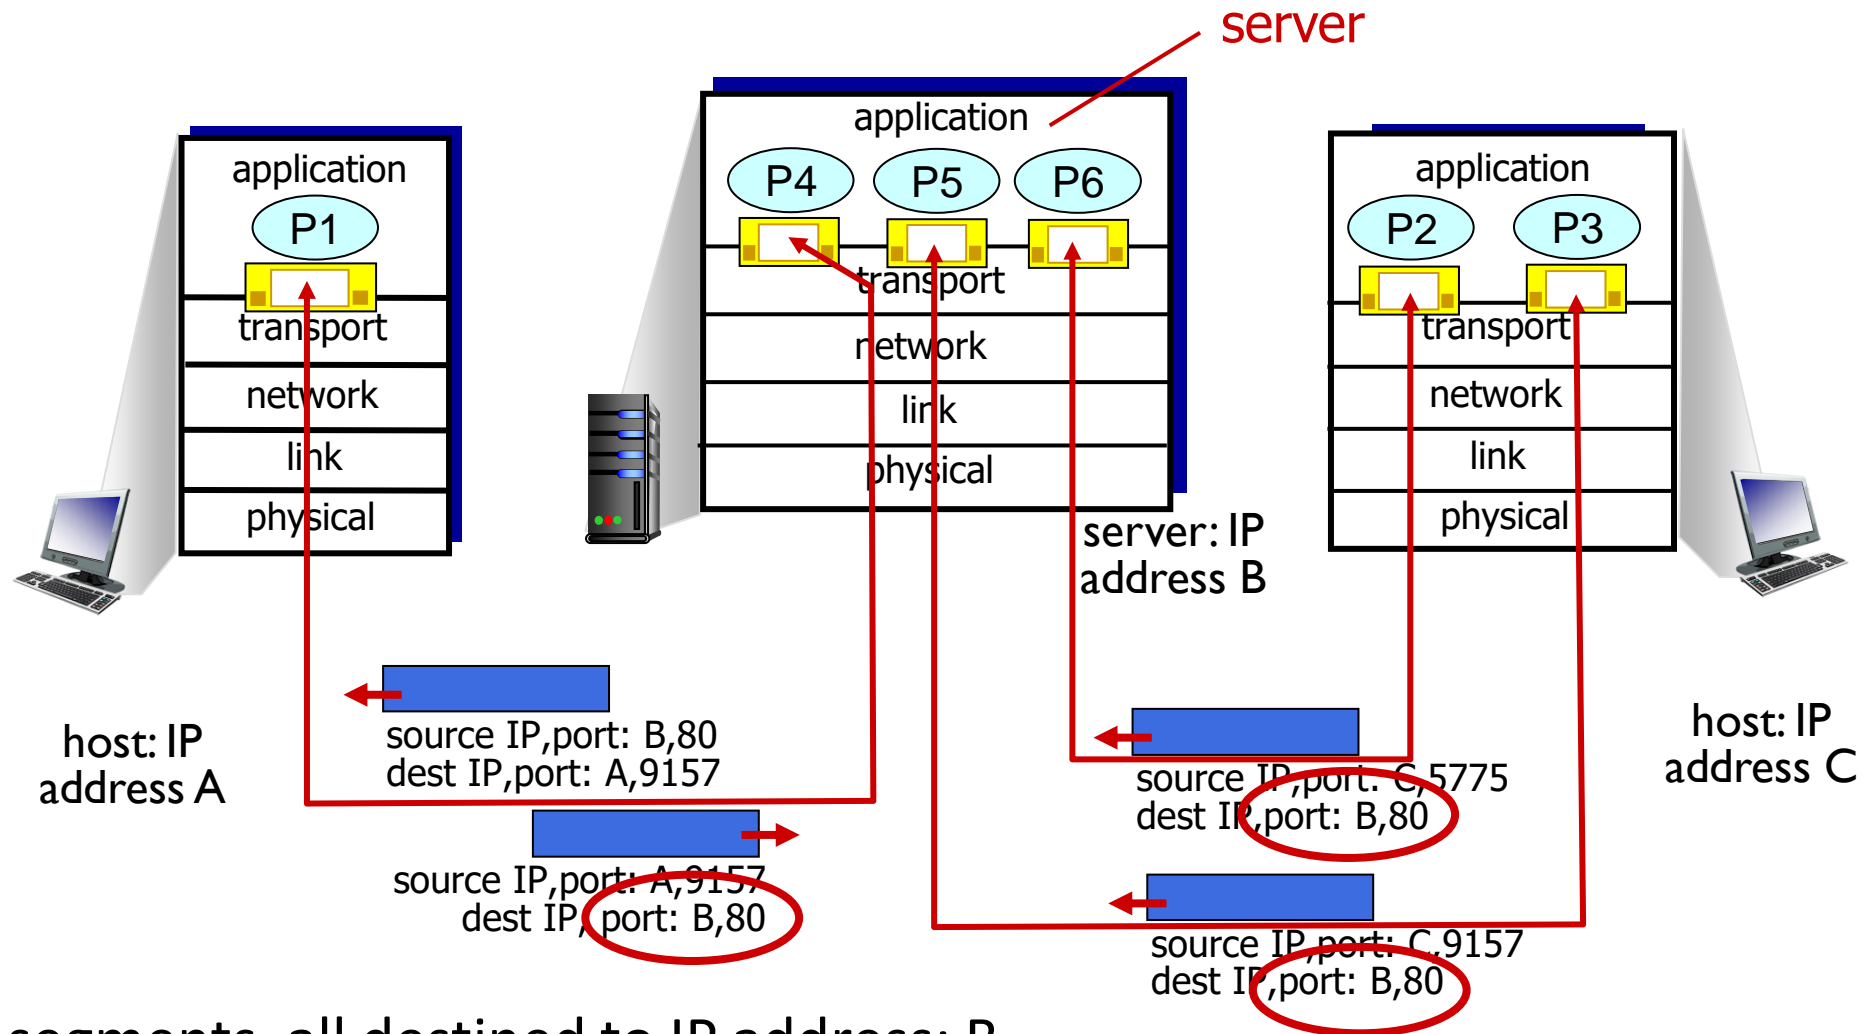
```
DatagramSocket mySocket1 =  
new DatagramSocket (5775);
```



Connection-oriented demultiplexing

- TCP socket identified by **4-tuple**:
 - *source IP address*
 - *source port number*
 - *dest IP address*
 - *dest port number*
- demux: receiver uses *all four values (4-tuple)* to direct segment to appropriate socket
- server may support many simultaneous TCP sockets:
 - each socket identified by its own 4-tuple
 - each socket associated with a different connecting client

Connection-oriented demultiplexing: example



Three segments, all destined to IP address: B,
dest port: 80 are demultiplexed to *different* sockets

Summary

- Multiplexing, demultiplexing: based on segment, datagram header field values
- **UDP:** demultiplexing using destination port number (only)
- **TCP:** demultiplexing using 4-tuple: source and destination IP addresses, and port numbers

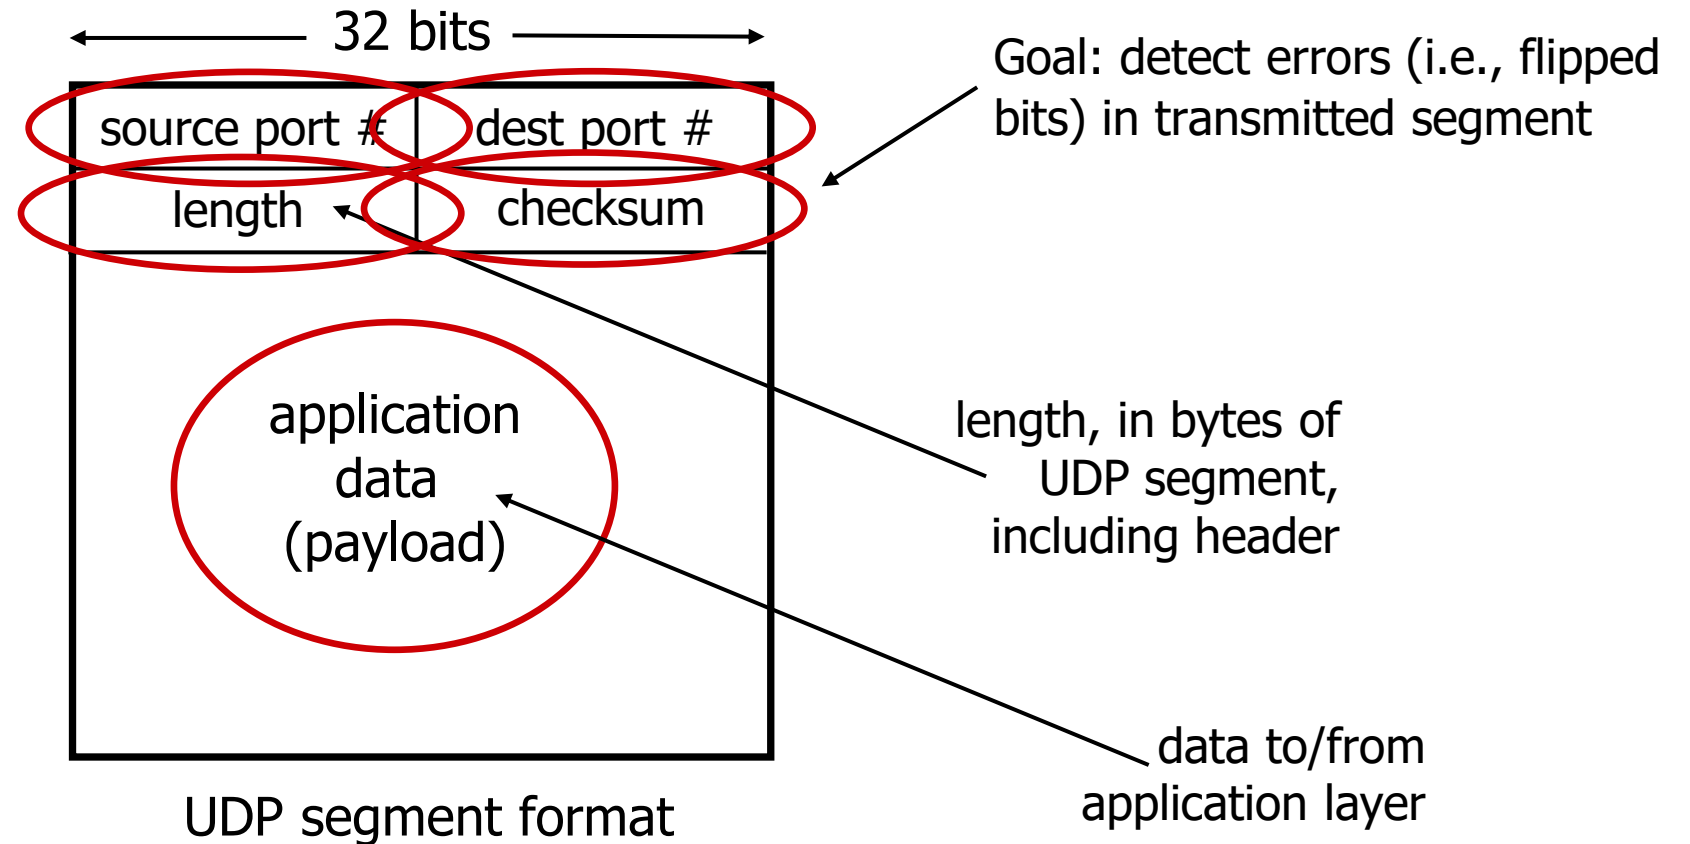
UDP: User Datagram Protocol

- “no frills,” “bare bones”
Internet transport protocol
- “best effort” service, UDP segments may be:
 - lost
 - delivered out-of-order to app
- *connectionless*:
 - no handshaking between UDP sender, receiver
 - each UDP segment handled independently of others

Why is there a UDP?

- no connection establishment (which can add delay)
- simple: no connection state at sender, receiver
- small header size
- no congestion control
 - UDP can blast away as fast as desired!
 - can function in the face of congestion

UDP segment header



Summary: UDP

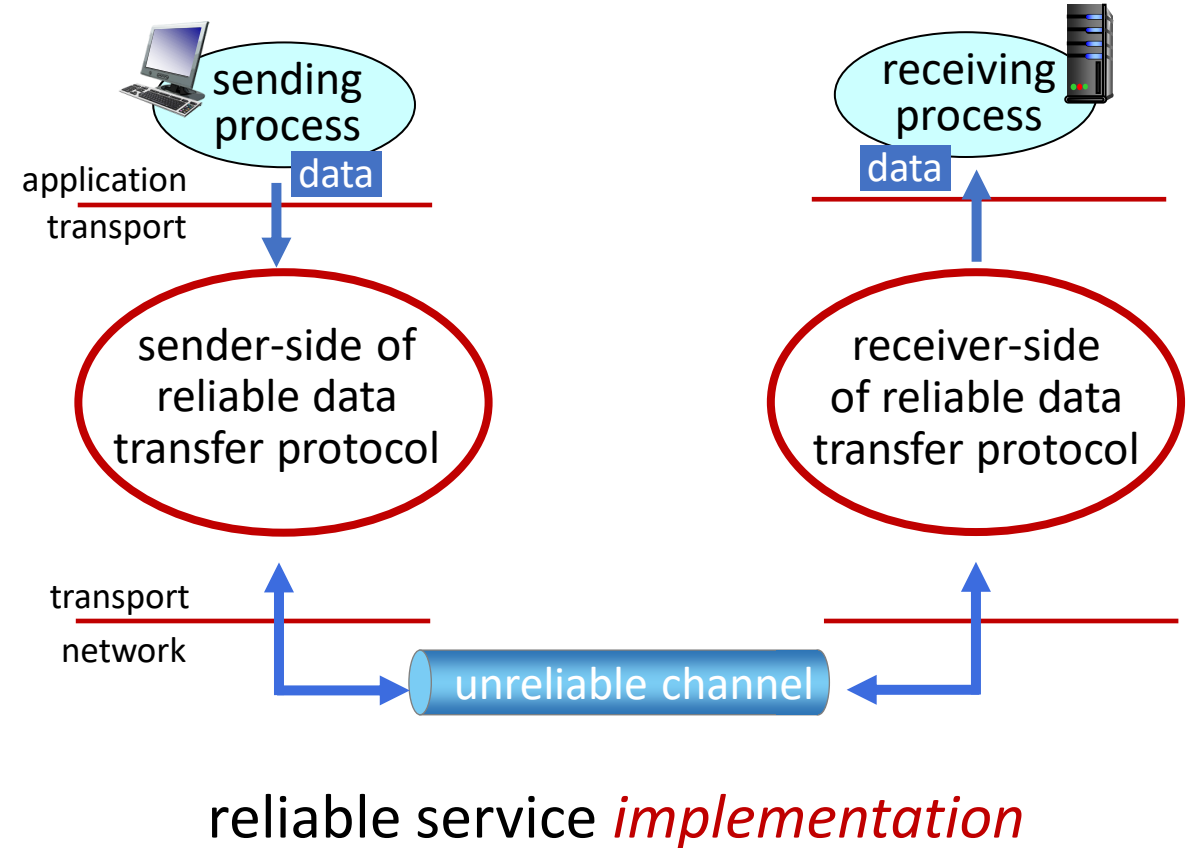

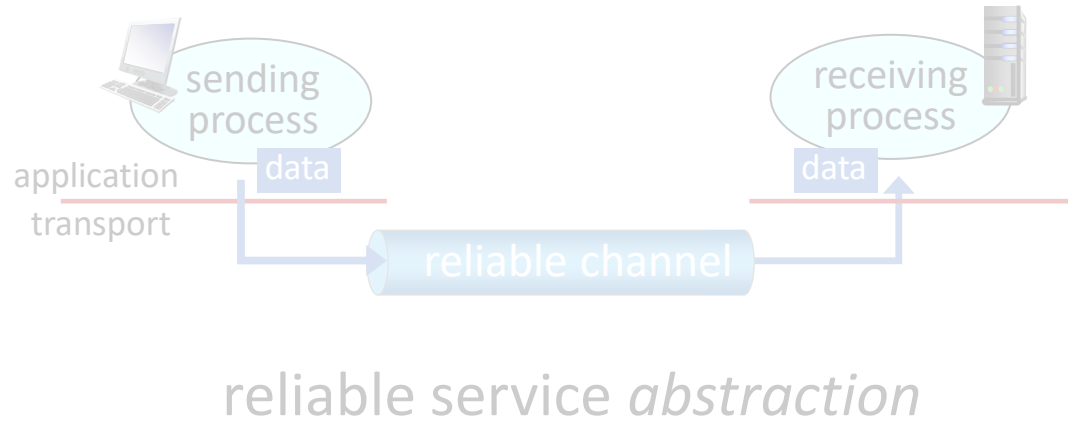
- “no frills” protocol:
 - segments may be lost, delivered out of order
 - best effort service: “send and hope for the best”
- UDP has its plusses:
 - no setup/handshaking needed (no delay incurred)
 - can function when network service is compromised
 - helps with reliability (checksum)
- build additional functionality on top of UDP in application layer

Principles of reliable data transfer



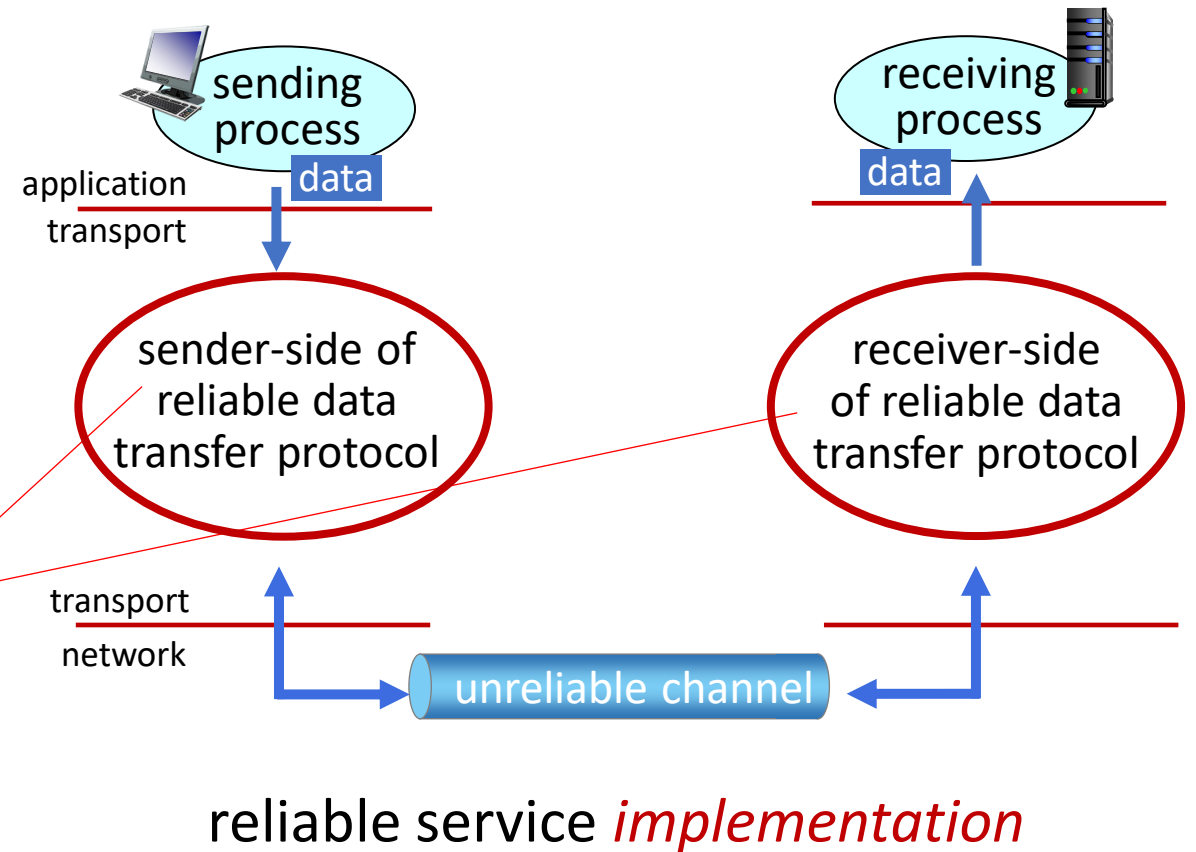
reliable service *abstraction*

Principles of reliable data transfer



Principles of reliable data transfer

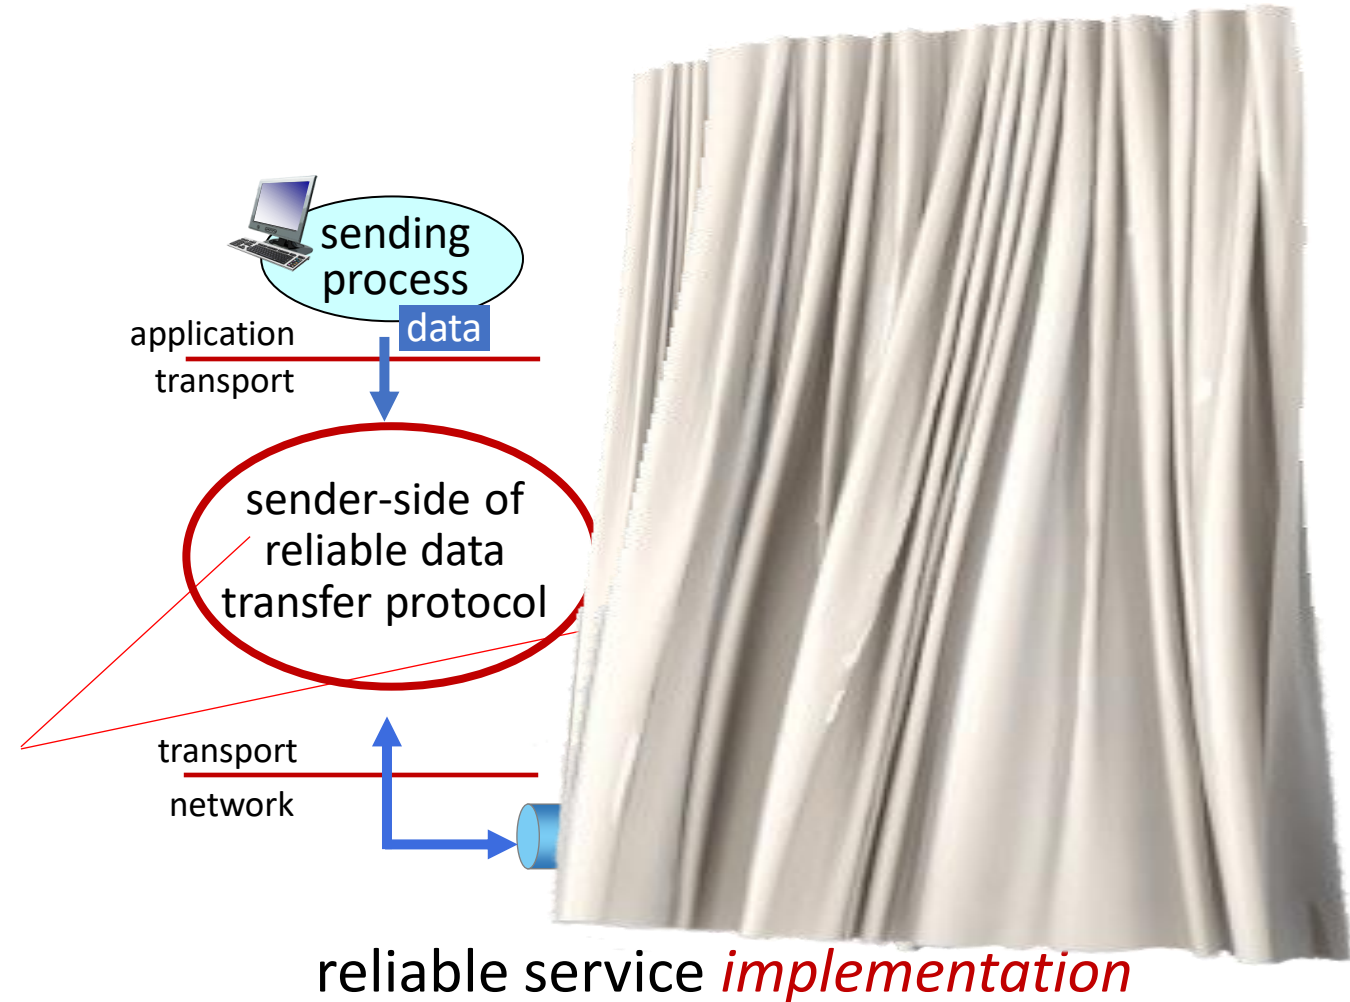
Complexity of reliable data transfer protocol will depend (strongly) on characteristics of unreliable channel (lose, corrupt, reorder data?)




Principles of reliable data transfer

Sender, receiver do *not* know the “state” of each other, e.g., was a message received?

- **unless communicated via a message**




Reliable data transfer



- Let's consider an «**ideal network**», meaning that it does not introduce
 - Bit errors
 - Discards of segments
 - Out-of-sequence segments
- The transport layer does not need to correct anything and the protocol is trivial
 - The sender sends segments in sequence (one after the other) and the receiver receives them without any need of further checks

Reliable data transfer




- Unfortunately, ideal networks do not exist...
- In a network with errors it is possible to introduce **positive acknowledgments** (ack) or **negative acknowledgments** (nack)
- Simple possible sender algorithm:


```

      IF ack
      THEN  $M_{i+1}$ 
      ELSE IF nack
           THEN  $M_i$ 
           ELSE ?
      
```
- However, **it does not work!** Ack/nack can also be subject to errors!
 - Possible mitigation strategy: if the ack/nack is corrupted, just retransmit the message → It does not work either! Impossible to understand at the receiver if the segment is a duplicate or not!

Reliable data transfer




- If the **segments are numbered**, there is no risk also in the case of duplicates

IF ack

THEN M_{i+1}

ELSE M_i \leftarrow A nack is received

Reliable data transfer

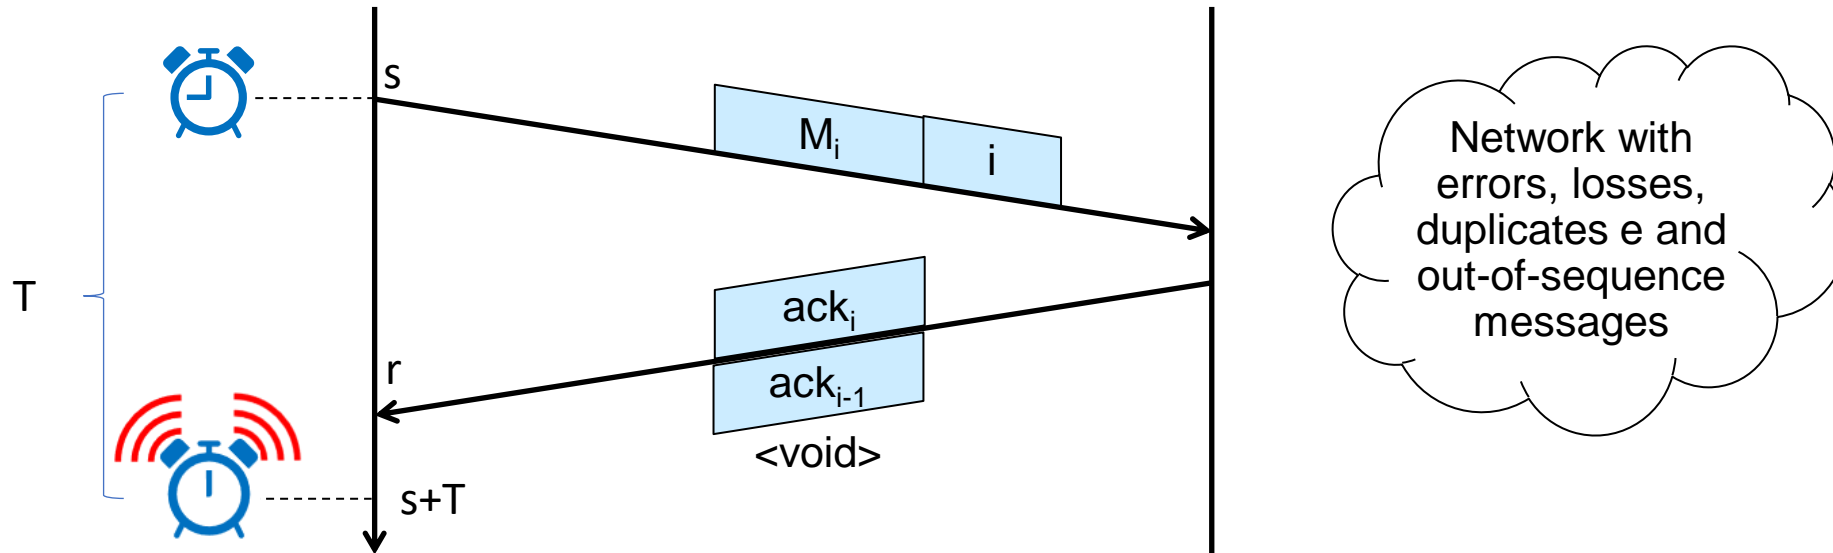


- By also **numbering the acks**, it is possible to avoid nacks thanks to the following rule: *a second ack_{i-1} is equal to a $nack_i$*

IF ack_i
 THEN M_{i+1}
 ELSE M_i \leftarrow An ack_{i-1} is received

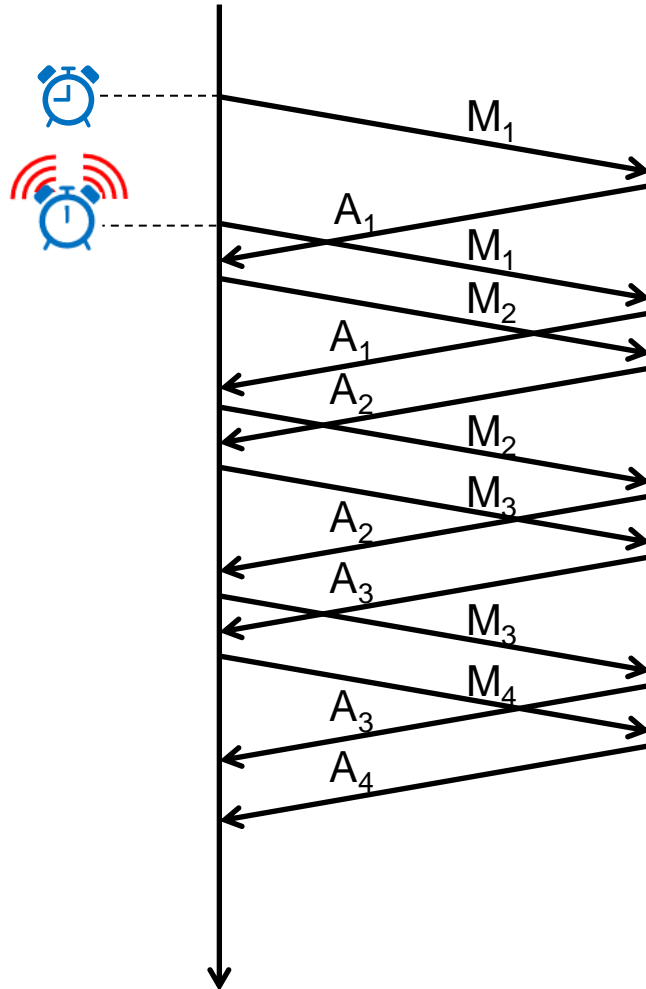
- Unfortunately, network with errors but without losses do not exist

Reliable data transfer




- By adding a **timer** we can handle losses (of segments or of acks)
 - IF $s+T$ is reached
 - THEN M_i , set $(s+T)+T$
 - ELSE IF $\text{ack}_i \leftarrow s+T$ has not been reached yet (r)
 - THEN M_{i+1} , set $r+T$
 - ELSE M_i , set $r+T$
- Note: what is shown is the case when $r \leq s+T$
 - In the case that $r \geq s+T$ the ack is ignored!
- This is the most sophisticated implementation of a **stop-and-wait** protocol

Reliable data transfer



- Setting the timer's deadline is a complex problem
- A **too short timer** may cause useless retransmissions
 - In the case seen so far (stop-and-wait) there may be even very long sequences of useless retransmissions
- A too long timer stops the transmission for long periods in the case of a loss

Reliable data transfer




- Transmitting a segment at a time and waiting for its ack before a further transmission (stop-and-wait) significantly limits performance

$$U = \frac{L/R}{RTT + L/R}$$


- Example: $RTT = 100\text{ ms}$, $L = 1\text{ kbyte}$,
 $R = 100\text{ Mbit/s} \rightarrow U = 0,008$

Reliable data transfer



- **Sliding window protocols** can transit up to W segments while waiting the ack of the first one

Reliable data transfer




- The condition for a continuous transmission is that the window will not close before the arrival of the first ack

$$W \cdot L/R \geq RTT + L/R$$

$$W \geq \frac{RTT \cdot R}{L} + 1$$


- In the example of two slides ago: $W \geq 1251$

Reliable data transfer



- What is retransmitted when a segment is lost?
- First possibility: **Go-Back-N** protocol
 - Transmit and slide the window after any ack reception
 - If $s+T$ is reached or repeated ack is received, retransmit all the segments from the last acknowledged segment
- There is no need of buffering at the receiver
 - Duplicate segments received after a loss are simply discarded
 - No further duplications once the segment in sequence is retransmitted
- Acks are cumulative
 - It is possible to recover from multiple ack losses

Reliable data transfer



- Second possibility: **Selective Repeat** protocol
 - Acks are individual
 - If $s+T$ is reached only the lost segment is retransmitted
- A buffer is needed at the receiver
 - Segments are reordered and sent to the application in the right order

Reliable data transfer

■ Observations:

- There are not only two possibilities (i.e., pure Go-Back-N and pure Selective Repeat): real protocols as TCP rely on hybrid solutions
- Sequence numbers of segments are represented with a finite number of bits
 - Some rules are defined to bind the maximum sequence number to the window's size, so that confusion is avoided when sequence numbers start back from zero:
 - Go-Back-N: $N \geq W + 1$
 - Selective Repeat: $N \geq 2W$